

```

[1]: import findspark
    findspark.init()
    #importing module to locate spark in VM

[2]: import pyspark
    from pyspark.sql.types import *
    #importing functions to run spark using python

[3]: sc = pyspark.SparkContext(appName="E10")
    #creates a spark context, using which we assigns the cluster of computer on which our code will run

    Setting default log level to "WARN".
    To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
    23/10/31 18:10:08 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

[4]: rdd1 = sc.textFile("/home/hduser/spark/nsedata.csv")
    #importing csv file as rdd

[5]: rdd1 = rdd1.filter(lambda x: "SYMBOL" not in x)
    #this is used to remove header row

[6]: rdd2 = rdd1.map(lambda x : x.split(","))
    #this is used to create a new rdd from the rdd1 by splitting each record

[7]: # Helper comment!: The goal is to find out the mean of the OPEN prices and the mean of the CLOSE price in one batch of tasks ...

[8]: rdd_open = rdd2.map(lambda x : (x[0]+"_open",float(x[2])))
    rdd_close = rdd2.map(lambda x : (x[0]+"_close",float(x[5])))
    #this creates two new rdds which contain the opening and closing prices
    #x[2] and x[5] are basically the positions containing the opening and closing prices

[3]: sc = pyspark.SparkContext(appName="E10")
    #creates a spark context, using which we assigns the cluster of computer on which our code will run

    Setting default log level to "WARN".
    To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
    23/10/31 18:10:08 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

[4]: rdd1 = sc.textFile("/home/hduser/spark/nsedata.csv")
    #importing csv file as rdd

[5]: rdd1 = rdd1.filter(lambda x: "SYMBOL" not in x)
    #this is used to remove header row

[6]: rdd2 = rdd1.map(lambda x : x.split(","))
    #this is used to create a new rdd from the rdd1 by splitting each record

[7]: # Helper comment!: The goal is to find out the mean of the OPEN prices and the mean of the CLOSE price in one batch of tasks ...

[8]: rdd_open = rdd2.map(lambda x : (x[0]+"_open",float(x[2])))
    rdd_close = rdd2.map(lambda x : (x[0]+"_close",float(x[5])))
    #this creates two new rdds which contain the opening and closing prices
    #x[2] and x[5] are basically the positions containing the opening and closing prices

[9]: rdd_united = rdd_open.union(rdd_close)
    #here we are combining both the rdds into a single rdd to do their analysis

[10]: reducedByKey = rdd_united.reduceByKey(lambda x,y: x+y)
    #to find the average we first need to find the sum of all
    #hence, here we are basically summing up the opening and closing prices for each entry

[11]: temp1 = rdd_united.map(lambda x: (x[0],1)).countByKey()
    countOfEachSymbol = sc.parallelize(temp1.items())
    #finding out the number of times a given entry occurred
    #we creates a new key value pair rdd
    #also we are counting the number of occurrence of each stock using the "countbykey" function
    #Parallelize is a method to create an RDD from an existing collection (For e.g Array) present in the driver.

[12]: symbol_sum_count = reducedByKey.join(countOfEachSymbol)
    #here we are combining the two rdds into a single one

[13]: averages = symbol_sum_count.map(lambda x : (x[0], x[1][0]/x[1][1]))
    # here we are finally finding the average of each of the symbols
    # so x[1][0]/x[1][1] is basically calculating sum of opening prices and closing prices of that particular symbol

[14]: averagesSorted = averages.sortByKey()
    #we are sorting the records present in the rdd based on symbol name

[15]: averagesSorted.saveAsTextFile("/home/hduser/spark/averages")
    # saving the results in folder named averages

[16]: sc.stop()

```

```

1 ('20MICRONS_close', 53.004122877930484)
2 ('20MICRONS_open', 53.32489894907032)
3 ('3IINFOTECH_close', 18.038803556992725)
4 ('3IINFOTECH_open', 18.17417138237672)
5 ('3MINDIA_close', 4520.343977364591)
6 ('3MINDIA_open', 4531.084518997574)
7 ('3RDROCK_close', 173.2137755102041)
8 ('3RDROCK_open', 173.18316326530612)
9 ('8KMILES_close', 480.73622047244095)
10 ('8KMILES_open', 481.63858267716535)
11 ('A2ZINFRA_close', 18.609433962264156)
12 ('A2ZINFRA_open', 18.73553459119497)
13 ('A2ZMES_close', 89.69389505549951)
14 ('A2ZMES_open', 90.46271442986883)
15 ('AANJANEYA_close', 441.84030249110316)
16 ('AANJANEYA_open', 440.93959074733095)
17 ('AARTIDRUGS_close', 312.94446240905427)
18 ('AARTIDRUGS_open', 312.832012934519)
19 ('AARTIIND_close', 127.70270816491507)
20 ('AARTIIND_open', 127.76463217461601)

```

```

In [2]: import pyspark
        from pyspark.sql.types import *

```

```

In [3]: sc = pyspark.SparkContext(appName="E10-01")

```

Setting default log level to "WARN".
 To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.
 23/10/28 17:12:24 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

```

In [4]: rdd7 = sc.textFile("/home/hduser/spark/nsedata.csv")

```

```

In [5]: rdd8 = rdd7.filter(lambda x: "SYMBOL" not in x)

```

```

In [6]: rdd9 = rdd8.map(lambda x : x.split(","))

```

```

In [17]: rdd_10 = rdd9.filter(lambda x: x[10] in "Apr" or "May" or "Jul" )

```

```

In [18]: rdd_high_1 = rdd_10.map(lambda x : (x[0]+"_high",float(x[3])))

```

```

In [19]: reducedByKey_1 = rdd_high_1.reduceByKey(lambda x,y: x+y)

```

```

In [20]: temp_1 = rdd_high_1.map(lambda x: (x[0],1)).countByKey()
        countOfEachSymbol_1 = sc.parallelize(temp_1.items())

```

```

In [21]: symbol_sum_count_1 = reducedByKey_1.join(countOfEachSymbol_1)

```

```

In [22]: averages_high = symbol_sum_count_1.map(lambda x : (x[0], x[1][0]/x[1][1]))

```

```

In [23]: averagesSorted = averages_high.sortByKey(ascending=False)

```

```

In [24]: averagesSorted.saveAsTextFile("/home/hduser/spark/averages_high_4")

```